

Timer.cc

```
int time = 40;
```

```
// dummy function because C++ does not allow pointers to member functions
```

```
static void TimerHandler(_int arg)
```

```
{ Timer *p = (Timer *)arg; p->TimerExpired(); }
```

```
//-----
```

```
// Timer::Timer
```

```
// Initialize a hardware timer device. Save the place to call
```

```
// on each interrupt, and then arrange for the timer to start
```

```
// generating interrupts.
```

```
//
```

```
// "timerHandler" is the interrupt handler for the timer device.
```

```
// It is called with interrupts disabled every time the
```

```
// the timer expires.
```

```
// "callArg" is the parameter to be passed to the interrupt handler.
```

```
// "doRandom" -- if true, arrange for the interrupts to occur
```

```
// at random, instead of fixed, intervals.
```

```
//-----
```

```
Timer::Timer(VoidFunctionPtr timerHandler, _int callArg, bool doRandom)
```

```
{
```

```
    randomize = doRandom;
```

```
    handler = timerHandler;
```

```
    arg = callArg;
```

```
    // schedule the first interrupt from the timer device
```

```
    interrupt->Schedule(TimerHandler, (_int) this, TimeOfNextInterrupt(),  
                        TimerInt);
```

```
}
```

```
//-----
```

```
// Timer::TimerExpired
```

```
// Routine to simulate the interrupt generated by the hardware
```

```
// timer device.
```

```
//-----
```

```
void
```

```
Timer::TimerExpired()
```

```
{
```

```
    /* Experiment 2 */
```

```
    /* Add code below to make the timer periodic. */
```

```
    interrupt->Schedule(TimerHandler, (_int) this, TimeOfNextInterrupt(),  
                        TimerInt);
```

```
    // invoke the Nachos interrupt handler for this device
```

```
    (*handler)(arg);
```

```
}
```

```
//-----  
// Timer::TimeOfNextInterrupt  
//   Return when the hardware timer device will next cause an interrupt.  
//   If randomize is turned on, make it a (pseudo-)random delay.  
//-----  
int  
Timer::TimeOfNextInterrupt()  
{  
    /* Experiment 2 */  
    /* Update below code so that it returns a fixed time quantum of 40 time ticks */  
  
    if (randomize)  
        return time;  
    else  
        return TimerTicks;  
}
```

Interrupt.cc

```
/* Experiment 2 */
```

```
/* Add code to remove pending interrupt from head of pending list. */
```

```
void
```

```
Interrupt::RemovePendingHead(){
```

```
    delete pending->Remove();
```

```
}
```

```
// String definitions for debugging messages
```

```
static char *intLevelNames[] = { "off", "on"};
```

```
static char *intTypeNames[] = { "timer", "disk", "console write",  
                                "console read", "network send", "network recv"};
```

```
//-----
```

```
// PendingInterrupt::PendingInterrupt
```

```
//     Initialize a hardware device interrupt that is to be scheduled
```

```
//     to occur in the near future.
```

```
//
```

```
//     "func" is the procedure to call when the interrupt occurs
```

```
//     "param" is the argument to pass to the procedure
```

```
//     "time" is when (in simulated time) the interrupt is to occur
```

```
//     "kind" is the hardware device that generated the interrupt
```

```
//-----
```

```
PendingInterrupt::PendingInterrupt(VoidFunctionPtr func, _int param, int time,
```

```
                                   IntType kind)
```

```
{
```

```
    handler = func;
```

```
    arg = param;
```

```
    when = time;
```

```
    type = kind;
```

```
}
```

```
//-----
```

```
// Interrupt::Interrupt
```

```
//     Initialize the simulation of hardware device interrupts.
```

```
//
```

```
//     Interrupts start disabled, with no interrupts pending, etc.
```

```
//-----
```

```
Interrupt::Interrupt()
```

```
{
```

```
    level = IntOff;
```

```
    pending = new List();
```

```
    inHandler = FALSE;
```

```
    yieldOnReturn = FALSE;
```

```
    status = SystemMode;
```

```
}
```

```

//-----
// Interrupt::~Interrupt
//      De-allocate the data structures needed by the interrupt simulation.
//-----
Interrupt::~Interrupt()
{
    while (!pending->IsEmpty())
        delete pending->Remove();
    delete pending;
}

//-----
// Interrupt::ChangeLevel
//      Change interrupts to be enabled or disabled, without advancing
//      the simulated time (normally, enabling interrupts advances the time).

//-----
// Interrupt::ChangeLevel
//      Change interrupts to be enabled or disabled, without advancing
//      the simulated time (normally, enabling interrupts advances the time).
//
//      Used internally.
//
//      "old" -- the old interrupt status
//      "now" -- the new interrupt status
//-----
void
Interrupt::ChangeLevel(IntStatus old, IntStatus now)
{
    level = now;
    DEBUG('i',"\\tinterrupts: %s -> %s\\n",intLevelNames[old],intLevelNames[now]);
}

```

```
//-----
// Interrupt::SetLevel
//      Change interrupts to be enabled or disabled, and if interrupts
//      are being enabled, advance simulated time by calling OneTick().
//
// Returns:
//      The old interrupt status.
// Parameters:
//      "now" -- the new interrupt status
//-----
IntStatus
Interrupt::SetLevel(IntStatus now)
{
    IntStatus old = level;

    ASSERT((now == IntOff) || (inHandler == FALSE)); // interrupt handlers are
                                                    // prohibited from enabling
                                                    // interrupts

    ChangeLevel(old, now); // change to new state
    if ((now == IntOn) && (old == IntOff))
        OneTick(); // advance simulated time
    return old;
}
```

```
//-----
// Interrupt::Enable
//      Turn interrupts on. Who cares what they used to be?
//      Used in ThreadRoot, to turn interrupts on when first starting up
//      a thread.
//-----
void
Interrupt::Enable()
{
    (void) SetLevel(IntOn);
}
```

```

//-----
// Interrupt::OneTick
//     Advance simulated time and check if there are any pending
//     interrupts to be called.
//
//     Two things can cause OneTick to be called:
//         interrupts are re-enabled
//         a user instruction is executed
//-----
void
Interrupt::OneTick()
{
    MachineStatus old = status;

// advance simulated time
    if (status == SystemMode) {
        stats->totalTicks += SystemTick;
        stats->systemTicks += SystemTick;
    } else {                                // USER_PROGRAM
        stats->totalTicks += UserTick;
        stats->userTicks += UserTick;
    }
    DEBUG('i', "\n== Tick %d ==\n", stats->totalTicks);

// check any pending interrupts are now ready to fire
    ChangeLevel(IntOn, IntOff);              // first, turn off interrupts
                                              // (interrupt handlers run with
                                              // interrupts disabled)
    while (CheckIfDue(FALSE))                // check for pending interrupts
        ;
    ChangeLevel(IntOff, IntOn);              // re-enable interrupts
    if (yieldOnReturn) {                     // if the timer device handler asked
                                              // for a context switch, ok to do it now
        yieldOnReturn = FALSE;
        status = SystemMode;                // yield is a kernel routine
        currentThread->Yield();
        status = old;
    }
}

```

```

//-----
// Interrupt::YieldOnReturn
//     Called from within an interrupt handler, to cause a context switch
//     (for example, on a time slice) in the interrupted thread,
//     when the handler returns.
//
//     We can't do the context switch here, because that would switch
//     out the interrupt handler, and we want to switch out the
//     interrupted thread.
//-----
void
Interrupt::YieldOnReturn()
{
    ASSERT(inHandler == TRUE);
    yieldOnReturn = TRUE;
}

//-----
// Interrupt::Idle
//     Routine called when there is nothing in the ready queue.
//
//     Since something has to be running in order to put a thread
//     on the ready queue, the only thing to do is to advance
//     simulated time until the next scheduled hardware interrupt.
//
//     If there are no pending interrupts, stop. There's nothing
//     more for us to do.
//-----
void
Interrupt::Idle()
{
    DEBUG('i', "Machine idling; checking for interrupts.\n");
    status = IdleMode;
    if (CheckIfDue(TRUE)) {           // check for any pending interrupts
        while (CheckIfDue(FALSE))    // check for any other pending
            ;                          // interrupts
        yieldOnReturn = FALSE;        // since there's nothing in the
                                     // ready queue, the yield is automatic
    }
    status = SystemMode;
    return;                           // return in case there's now
                                     // a runnable thread
}
// if there are no pending interrupts, and nothing is on the ready
// queue, it is time to stop. If the console or the network is
// operating, there are *always* pending interrupts, so this code
// is not reached. Instead, the halt must be invoked by the user program.
DEBUG('i', "Machine idle. No interrupts to do.\n");
printf("No threads ready or runnable, and no pending interrupts.\n");
printf("Assuming the program completed.\n");
Halt();
}

```

```
//-----
// Interrupt::Halt
//     Shut down Nachos cleanly, printing out performance statistics.
//-----
void
Interrupt::Halt()
{
    printf("Machine halting!\n\n");
    stats->Print();
    Cleanup();    // Never returns.
}
```

```
//-----
// Interrupt::Schedule
//     Arrange for the CPU to be interrupted when simulated time
//     reaches "now + when".
//
//     Implementation: just put it on a sorted list.
//
//     NOTE: the Nachos kernel should not call this routine directly.
//     Instead, it is only called by the hardware device simulators.
//
//     "handler" is the procedure to call when the interrupt occurs
//     "arg" is the argument to pass to the procedure
//     "fromNow" is how far in the future (in simulated time) the
//             interrupt is to occur
//     "type" is the hardware device that generated the interrupt
//-----
void
Interrupt::Schedule(VoidFunctionPtr handler, _int arg, int fromNow, IntType type)
{
    int when = stats->totalTicks + fromNow;
    PendingInterrupt *toOccur = new PendingInterrupt(handler, arg, when, type);

    DEBUG('i', "Scheduling interrupt handler the %s at time = %d\n",
          intTypeNames[type], when);
    ASSERT(fromNow > 0);

    pending->SortedInsert(toOccur, when);
}
```



```

//-----
// Interrupt::CheckIfDue
//     Check if an interrupt is scheduled to occur, and if so, fire it off.
//
// Returns:
//     TRUE, if we fired off any interrupt handlers
// Params:
//     "advanceClock" -- if TRUE, there is nothing in the ready queue,
//                     so we should simply advance the clock to when the next
//                     pending interrupt would occur (if any). If the pending
//                     interrupt is just the time-slice daemon, however, then
//                     we're done!
//-----
bool
Interrupt::CheckIfDue(bool advanceClock)
{
    MachineStatus old = status;
    int when;
    ASSERT(level == IntOff);           // interrupts need to be disabled,
                                      // to invoke an interrupt handler

    if (DebugIsEnabled('i'))
        DumpState();
    PendingInterrupt *toOccur =
        (PendingInterrupt *)pending->SortedRemove(&when);

    if (toOccur == NULL)               // no pending interrupts
        return FALSE;

    if (advanceClock && when > stats->totalTicks) { // advance the clock
        stats->idleTicks += (when - stats->totalTicks);
        stats->totalTicks = when;
    } else if (when > stats->totalTicks) { // not time yet, put it back
        pending->SortedInsert(toOccur, when);
        return FALSE;
    }
}
// Check if there is nothing more to do, and if so, quit
if ((status == IdleMode) && (toOccur->type == TimerInt)
    && pending->IsEmpty()) {
    pending->SortedInsert(toOccur, when);
    return FALSE;
}
DEBUG('i', "Invoking interrupt handler for the %s at time %d\n",
      intTypeNames[toOccur->type], toOccur->when);
#ifdef USER_PROGRAM
    if (machine != NULL)
        machine->DelayedLoad(0, 0);
#endif
inHandler = TRUE;
status = SystemMode;                // whatever we were doing,
                                    // we are now going to be
                                    // running in the kernel
(*(toOccur->handler))(toOccur->arg); // call the interrupt handler
status = old;                       // restore the machine status
inHandler = FALSE;

```

```

    delete toOccur;
    return TRUE;
}

//-----
// PrintPending
//     Print information about an interrupt that is scheduled to occur.
//     When, where, why, etc.
//-----
static void
PrintPending(_int arg)
{
    PendingInterrupt *pend = (PendingInterrupt *)arg;

    printf("Interrupt handler %s, scheduled at %d\n",
        intTypeNames[pend->type], pend->when);
}

//-----
// DumpState
//     Print the complete interrupt state - the status, and all interrupts
//     that are scheduled to occur in the future.
//-----
void
Interrupt::DumpState()
{
    printf("Time: %d, interrupts %s\n", stats->totalTicks,
        intLevelNames[level]);

    printf("Pending interrupts:\n");
    fflush(stdout);
    pending->Mapcar(PrintPending);
    printf("End of pending interrupts\n");
    fflush(stdout);
}

```

System.cc

```
//-----
// TimerInterruptHandler
//     Interrupt handler for the timer device. The timer device is
//     set up to interrupt the CPU periodically (once every TimerTicks).
//     This routine is called each time there is a timer interrupt,
//     with interrupts disabled.
//
//     Note that instead of calling Yield() directly (which would
//     suspend the interrupt handler, not the interrupted thread
//     which is what we wanted to context switch), we set a flag
//     so that once the interrupt handler is done, it will appear as
//     if the interrupted thread called Yield at the point it is
//     was interrupted.
//
//     "dummy" is because every interrupt handler takes one argument,
//     whether it needs it or not.
//-----
static void
TimerInterruptHandler(int dummy)
{
    if (interrupt->getStatus() != IdleMode)
        interrupt->YieldOnReturn();
}

//-----
// Initialize
//     Initialize Nachos global data structures. Interpret command
//     line arguments in order to determine flags for the initialization.
//
//     "argc" is the number of command line arguments (including the name
//     of the command) -- ex: "nachos -d +" -> argc = 3
//     "argv" is an array of strings, one for each command line argument
//     ex: "nachos -d +" -> argv = {"nachos", "-d", "+"}
//-----
void
Initialize(int argc, char **argv)
{
    /* Experiment 2 */
    /*Identify where the timer is initialized in this file. Activate the initialization of the timer by updating
    appropriate variables.*/

    int argCount;
    char* debugArgs = "";
    bool randomYield = TRUE;

#ifdef USER_PROGRAM
    bool debugUserProg = FALSE;    // single step user program
#endif
#ifdef FILESYS_NEEDED
    bool format = FALSE;    // format disk
#endif
#endif
```

```

#ifdef NETWORK
    double rely = 1;           // network reliability
    int netname = 0;           // UNIX socket name
#endif

for (argc--, argv++; argc > 0; argc -= argCount, argv += argCount) {
    argCount = 1;
    if (!strcmp(*argv, "-d")) {
        if (argc == 1)
            debugArgs = "+";    // turn on all debug flags
        else {
            debugArgs = *(argv + 1);
            argCount = 2;
        }
    } else
    {
        ASSERT(argc > 1);
        RandomInit(10);         // initialize pseudo-random
                                // number generator w/ 10

        randomYield = TRUE;
        argCount = 2;
    }
}

#ifdef USER_PROGRAM
    if (!strcmp(*argv, "-s"))
        debugUserProg = TRUE;
#endif

#ifdef FILESYS_NEEDED
    if (!strcmp(*argv, "-f"))
        format = TRUE;
#endif

#ifdef NETWORK
    if (!strcmp(*argv, "-l")) {
        ASSERT(argc > 1);
        rely = atof(*(argv + 1));
        argCount = 2;
    } else if (!strcmp(*argv, "-m")) {
        ASSERT(argc > 1);
        netname = atoi(*(argv + 1));
        argCount = 2;
    }
}

#endif

    DebugInit(debugArgs);           // initialize DEBUG messages
    stats = new Statistics();        // collect statistics
    interrupt = new Interrupt();     // start up interrupt handling
    scheduler = new Scheduler();     // initialize the ready queue
    if (randomYield) {               // start the timer (if needed)
        timer = new Timer(TimerInterruptHandler, 0, randomYield);
        /* Experiment 2 */
        /* Debug message to denote scheduling of timer interrupt */
        DEBUG('i', "**** Timer interrupt scheduled at %d\n", stats->totalTicks+timer-
>TimeOfNextInterrupt());
    }

```

```

threadToBeDestroyed = NULL;

// We didn't explicitly allocate the current thread we are running in.
// But if it ever tries to give up the CPU, we better have a Thread
// object to save its state.
currentThread = new Thread("main");
currentThread->setStatus(RUNNING);

interrupt->Enable();
CallOnUserAbort(Cleanup);           // if user hits ctrl-C

#ifdef USER_PROGRAM
    machine = new Machine(debugUserProg); // this must come first
#endif
#ifdef CHANGED
    processes = new SList;
    processes->Add(new Process, 0, currentThread);
#endif
    console = new Console(0, 0, ReadConsoleAvail, WriteConsoleDone, 0);
    synchConsole = new SynchConsole;
#endif

#ifdef FILESYS
    synchDisk = new SynchDisk("DISK");
#endif

#ifdef FILESYS_NEEDED
    fileSystem = new FileSystem(format);
#endif

#ifdef NETWORK
    postOffice = new PostOffice(netname, rely, 10);
#endif

#ifdef CHANGED
    memoryTable = new MemoryTable[NumPhysPages];
#endif
}

//-----
// Cleanup
//     Nachos is halting. De-allocate global data structures.
//-----
void
Cleanup()
{
    printf("\nCleaning up...\n");
#ifdef NETWORK
    delete postOffice;
#endif
}

```

```
#ifdef USER_PROGRAM
    delete machine;
#endif
```

```
#ifdef FILESYS_NEEDED
    delete fileSystem;
#endif
```

```
#ifdef FILESYS
    delete synchDisk;
#endif
```

```
    delete timer;
```

```
#ifdef CHANGED
#ifdef USER_PROGRAM
    delete processes;
    delete console;
    delete synchConsole;
#endif
#endif
```

```
    delete scheduler;
    delete interrupt;
```

```
#ifdef CHANGED
    delete [] memoryTable;
#endif
```

```
    Exit(0);
}
```

Thread.cc

```
//-----
// Thread::Thread
//     Initialize a thread control block, so that we can then call
//     Thread::Fork.
//
//     "threadName" is an arbitrary string, useful for debugging.
//-----
Thread::Thread(char* threadName)
{
    name = threadName;
    stackTop = NULL;
    stack = NULL;
    status = JUST_CREATED;
#ifdef CHANGED
    join_thereP = new Semaphore("Join?", 0); // create Finish's sem. to
                                           // see if Join is there
#endif
#ifdef USER_PROGRAM
    space = NULL;
#endif
}

//-----
// Thread::~~Thread
//     De-allocate a thread.
//
//     NOTE: the current thread *cannot* delete itself directly,
//     since it is still running on the stack that we need to delete.
//
//     NOTE: if this is the main thread, we can't delete the stack
//     because we didn't allocate it -- we got it automatically
//     as part of starting up Nachos.
//-----
Thread::~~Thread()
{
    DEBUG('t', "Deleting thread %s #%i\n", name, pid);

    ASSERT(this != currentThread);
#ifdef CHANGED
    delete join_thereP;
#endif
    if (stack != NULL)
        DeallocBoundedArray((char *) stack, StackSize * sizeof(int));
}
```

```
//-----
// Thread::Fork
//     Invoke (*func)(arg), allowing caller and callee to execute
//     concurrently.
//
//     NOTE: although our definition allows only a single integer argument
//     to be passed to the procedure, it is possible to pass multiple
//     arguments by making them fields of a structure, and passing a pointer
//     to the structure as "arg".
//
//     Implemented as the following steps:
//         1. Allocate a stack
//         2. Initialize the stack so that a call to SWITCH will
//         cause it to run the procedure
//         3. Put the thread on the ready queue
//
//     "func" is the procedure to run concurrently.
//     "arg" is a single argument to be passed to the procedure.
//     "joinP" is 0 if this thread cannot be joined and
//         1 if this thread will be joined
//-----
#ifdef CHANGED
void
Thread::Fork(VoidFunctionPtr func, int arg, int joinP)
{
    DEBUG('t', "Forking thread %s #%i with func=0x%x, arg=%d, join=%s\n",
        name, pid, (int) func, arg, (joinP ? "YES" : "NO"));

    will_joinP = joinP;                // remember if you are joined for
                                        // the finish procedure

    StackAllocate(func, arg);

    IntStatus oldLevel = interrupt->SetLevel(IntOff);
    scheduler->ReadyToRun(this);        // ReadyToRun assumes that interrupts
                                        // are disabled!
    (void) interrupt->SetLevel(oldLevel);
}

```



```

//-----
// Thread::Join
//     Wait until the specified thread has completed running before
//     the code from the point of the Join call on is run. Both the
//     Join procedure and the thread will return when, and only when,
//     both have finished.
//-----
void Thread::Join(Thread *forked) {
    DEBUG('j', "Thread %s #%i is calling join on thread %s #%i\n", name, pid,
        forked->getName(), forked->pid);

    Semaphore *join_sem = new Semaphore("in_join", 0); // Sem. to see if
                                                    // Finish arrived
    forked->JoinHit(join_sem); // call in forked to send ^^ sem.
                            // and tell Finish's sem. that
                            // you are here
    join_sem->P(); // wait for Finish
    delete join_sem;
}

//-----
// Thread::JoinHit
//     Allows the function Join to access the semaphore in the thread's
//     class that is being joined, in order to let it know that Join
//     has been called. Also it sends it's semaphore so it can know
//     if the joined thread has gotten to it's finish state yet.
//-----
void Thread::JoinHit(class Semaphore *caller) {
    DEBUG('j', "In JoinHit for Thread %s #%i.\n", getName(), pid);

    join_wait = caller; // Get Join's sem. for Finish to V()
    join_thereP->V(); // Let Finish know Join is here
}
#endif

```

```
//-----
// Thread::CheckOverflow
//     Check a thread's stack to see if it has overrun the space
//     that has been allocated for it. If we had a smarter compiler,
//     we wouldn't need to worry about this, but we don't.
//
//     NOTE: Nachos will not catch all stack overflow conditions.
//     In other words, your program may still crash because of an overflow.
//
//     If you get bizarre results (such as seg faults where there is no code)
//     then you *may* need to increase the stack size. You can avoid stack
//     overflows by not putting large data structures on the stack.
//     Don't do this: void foo() { int bigArray[10000]; ... }
//-----
void
Thread::CheckOverflow()
{
    if (stack != NULL)
#ifdef HOST_SNAKE                // Stacks grow upward on the Snakes
        ASSERT(stack[StackSize - 1] == STACK_FENCEPOST);
#else
        ASSERT(*stack == STACK_FENCEPOST);
#endif
}
```

[illegible]

```

(void) interrupt->SetLevel(IntOff);
ASSERT(this == currentThread);

DEBUG('t', "Finishing thread %s #i\n", getName(), pid);

threadToBeDestroyed = currentThread;

/* Experiment 2 */
/* Add code here to reset the timer interrupt so that the next
   interrupt is triggered after 40 time ticks from now.
*/

interrupt -> ScheduHandler, (_int) this, timer ->TimeOfNextInterrupt(), TimerInt);
interrupt -> RemovePendingHead();

    Sleep();                                // invokes SWITCH
}
// not reached
}
#endif

//-----
// Thread::Yield
// Relinquish the CPU if any other thread is ready to run.
// If so, put the thread on the end of the ready list, so that
// it will eventually be re-scheduled.
//
// NOTE: returns immediately if no other thread on the ready queue.
// Otherwise returns when the thread eventually works its way
// to the front of the ready list and gets re-scheduled.
//
// NOTE: we disable interrupts, so that looking at the thread
// on the front of the ready list, and switching to it, can be done
// atomically. On return, we re-set the interrupt level to its
// original state, in case we are called with interrupts disabled.
//
// Similar to Thread::Sleep(), but a little different.
//-----
void
Thread::Yield ()
{
    Thread *nextThread;
    IntStatus oldLevel = interrupt->SetLevel(IntOff);

    ASSERT(this == currentThread);

    DEBUG('t', "Yielding thread %s #i\n", getName(), pid);

    nextThread = scheduler->FindNextToRun();
    if (nextThread != NULL) {
        scheduler->ReadyToRun(this);
        scheduler->Run(nextThread);
    }
}

```

```

    }
    (void) interrupt->SetLevel(oldLevel);
}

```

```

//-----
// Thread::Sleep
//     Relinquish the CPU, because the current thread is blocked
//     waiting on a synchronization variable (Semaphore, Lock, or Condition).
//     Eventually, some thread will wake this thread up, and put it
//     back on the ready queue, so that it can be re-scheduled.
//
//     NOTE: if there are no threads on the ready queue, that means
//     we have no thread to run. "Interrupt::Idle" is called
//     to signify that we should idle the CPU until the next I/O interrupt
//     occurs (the only thing that could cause a thread to become
//     ready to run).
//
//     NOTE: we assume interrupts are already disabled, because it
//     is called from the synchronization routines which must
//     disable interrupts for atomicity. We need interrupts off
//     so that there can't be a time slice between pulling the first thread
//     off the ready list, and switching to it.
//-----

```

```

void
Thread::Sleep ()
{
    Thread *nextThread;

    ASSERT(this == currentThread);
    ASSERT(interrupt->getLevel() == IntOff);

    DEBUG('t', "Sleeping thread %s #%i\n", getName(), pid);

    status = BLOCKED;
    while ((nextThread = scheduler->FindNextToRun()) == NULL)
        interrupt->Idle();    // no one to run, wait for an interrupt

    scheduler->Run(nextThread); // returns when we've been signalled
}

```

```

//-----

```

```
// ThreadFinish, InterruptEnable, ThreadPrint
//      Dummy functions because C++ does not allow a pointer to a member
//      function. So in order to do this, we create a dummy C function
//      (which we can pass a pointer to), that then simply calls the
//      member function.
//-----
```

```
static void ThreadFinish() { currentThread->Finish(); }
static void InterruptEnable() { interrupt->Enable(); }
void ThreadPrint(int arg){ Thread *t = (Thread *)arg; t->Print(); }
```

```
//-----
// Thread::StackAllocate
//      Allocate and initialize an execution stack. The stack is
//      initialized with an initial stack frame for ThreadRoot, which:
//          enables interrupts
//          calls (*func)(arg)
//          calls Thread::Finish
//
//      "func" is the procedure to be forked
//      "arg" is the parameter to be passed to the procedure
//-----
```

```
void
Thread::StackAllocate (VoidFunctionPtr func, int arg)
{
    stack = (int *) AllocBoundedArray(StackSize * sizeof(int));
```

```
#ifdef HOST_SNAKE
    // HP stack works from low addresses to high addresses
    stackTop = stack + 16;    // HP requires 64-byte frame marker
    stack[StackSize - 1] = STACK_FENCEPOST;
#else
    // i386 & MIPS & SPARC stack works from high addresses to low addresses
#endif
#ifdef HOST_SPARC
    // SPARC stack must contains at least 1 activation record to start with.
    stackTop = stack + StackSize - 96;
#else // HOST_MIPS || HOST_i386
    stackTop = stack + StackSize - 4;    // -4 to be on the safe side!
#endif
#ifdef HOST_i386
    // the 80386 passes the return address on the stack. In order for
    // SWITCH() to go to ThreadRoot when we switch to this thread, the
    // return address used in SWITCH() must be the starting address of
    // ThreadRoot.
    * (--stackTop) = (int)ThreadRoot;
#endif
#endif // HOST_SPARC
    *stack = STACK_FENCEPOST;
#endif // HOST_SNAKE
```

```
machineState[PCState] = (int) ThreadRoot;
machineState[StartupPCState] = (int) InterruptEnable;
machineState[InitialPCState] = (int) func;
```

```

    machineState[InitialArgState] = arg;
    machineState[WhenDonePCState] = (int) ThreadFinish;
}

#ifdef USER_PROGRAM
#include "machine.h"

//-----
// Thread::SaveUserState
//     Save the CPU state of a user program on a context switch.
//
//     Note that a user program thread has *two* sets of CPU registers --
//     one for its state while executing user code, one for its state
//     while executing kernel code. This routine saves the former.
//-----
void
Thread::SaveUserState()
{
    for (int i = 0; i < NumTotalRegs; i++)
        userRegisters[i] = machine->ReadRegister(i);
}

//-----
// Thread::RestoreUserState
//     Restore the CPU state of a user program on a context switch.
//
//     Note that a user program thread has *two* sets of CPU registers --
//     one for its state while executing user code, one for its state
//     while executing kernel code. This routine restores the former.
//-----
void
Thread::RestoreUserState()
{
    for (int i = 0; i < NumTotalRegs; i++)
        machine->WriteRegister(i, userRegisters[i]);
}
#endif

```

Interrupt.h

```
void RemovePendingHead(); //add this line
```

