

1)

a)

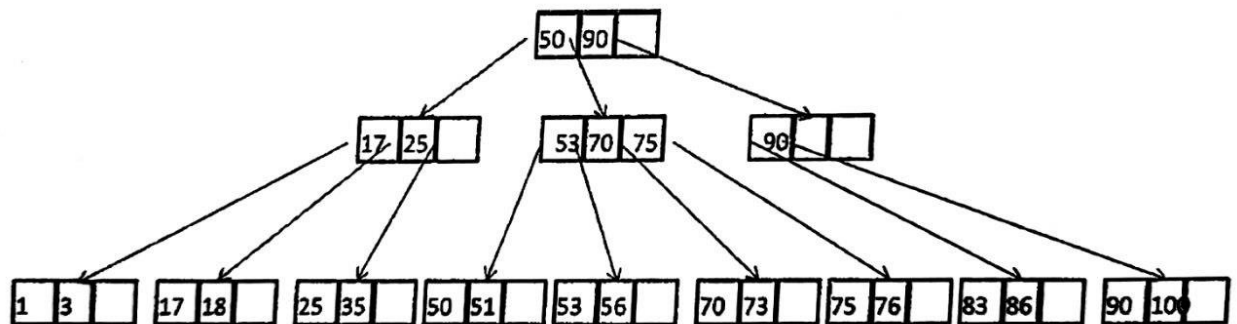
- i) Transfer time
- ii) Query time (insert, update, delete)
- iii) A because more records can fit into one block
- iv) Yes, because it has one index entry for each data record

b) Since prefetching algorithm is used,


$$\begin{aligned} \text{I/O cost} &= R + P + (\# \text{ of blocks} - 1) * \max(R, P) \\ &= R + 4P \quad [\text{since } R < P] \end{aligned}$$

c)


i)



ii)



iii)



The bold arrows indicates the search path.

d)  $(19+1)^{h-1} \cdot 19 \geq 1000$

$h \geq 3$

Hence, the minimum height of the B+ Tree is 3

e) For data, need  $100000/10 = 10000$  blocks

For index,

$79^{h-1} \cdot 80 \geq 100000$

$h = 3$  ( $h$  is the height of B+ tree)


Need  $1+80+80^2 = 6481$  blocks

Totally need 16481 blocks


2)

a)

i)



ii)



b)  $32 \cdot 10 \cdot 5 = 1600$  tuples

c) Yes.

Require less index space.

Update operations (insert, delete) are easier and more efficient.

d) B+ tree index (take advantage of range query) on the attribute 'a', clustered (take advantage of range query)

B+ tree index (take advantage of range query) on the attribute 'c', clustered (take advantage of range query)

e)  $B(R1) = 100000/200 = 500$  blocks

$B(R2) = 50000/500 = 100$  blocks

$M = 101$

Using block-based nested loop join

$\text{minCost} = B(R2) + B(R1) \cdot \text{upperBound}(B(R2)/(M-1)) = 600$  block access

3)

a)

i)  $B(R) = 1500, B(S) = 460, M = 101$

$B(R) + B(S) < M^2 \rightarrow$  Refined sort-merge join algorithm can be applied

I/O cost =  $3 * (B(R) + B(S)) = 5880$  disk I/O

ii) Read 1500 block of R, sort, write back 150 sorted sublists  $\rightarrow$  cost  $1500 * 2 = 3000$  I/O

Read 460 block of R, sort, write back 4 sorted sublists, the last sublist contains 60 last block will be kept in memory  $\rightarrow$  cost  $460 + 400 = 860$  I/O

Merge phase: read blocks in sublists of R and S linearly  $\rightarrow$  cost  $1500 + 400 = 1900$  I/O

Total cost  $3000 + 860 + 1900 = 5760$  I/O

Number of I/O(s) saved:  $5880 - 5760 = 120$  I/O

Actually the number of I/O saved is  $2 * \text{size of sublist kept in memory}$

iii) Cost of index-based join:  $B(S) + T(S) * T(R) / V(R, 'x') = 460 + 9200 * 15000 / 50$

Cost of hash join:  $3(B(S) + B(R)) = 3 * (460 + 1500)$

Hence index-based join is a much worse choice in the case.

iv)  $T(R) = 15000, T(S) = 9200$

Number of tuples satisfy  $y > 5$ :  $T(R) / 3 = 5000 \rightarrow X$

Number of tuples satisfy  $z \geq 9$ :  $0.2 * T(S) = 1840 \rightarrow Y$

(0.2 since  $1 \leq z \leq 10$  and  $z$  values are evenly distributed in S)

Join of X and Y:  $T(X) * T(Y) / \max(V(X, 'x'), V(Y, 'x')) = 5000 * 1840 / 50 = 184,000$

b) **Materialization**

Order the operations (an operation is not performed until the argument(s) below have been performed)

Store the result of each operation on disk until it is needed by another operation

**Pipelining**

Interleave the execution of several operations

The tuples produced by one operation are passed directly to the operation that it

No need to store intermediate relation in disk

Several operations must share main memory at any time

4)

a) **Atomicity:** A transaction is either performed in its entirety (all its steps) or not performed at all

**Correctness:** If T starts with DB in consistent state then T executes in isolation, T leaves consistent state

**Isolation:** A transaction should appear as though it is executed in isolation from other transactions

**Durability:** Changes applied to the database by a committed transaction must persist in the database

b)

i) T1 does not need to be redone because it is committed in LogID4 and successfully written to disk after <END CKPT>

- T2 needs to be redone because it's uncommitted (its commit log is LogID10 which is after the crash point)
- ii) T1 is done hence (T1, A, 2) is updated on disk  
If T2 successfully write to disk before the crash point  $\rightarrow (A, B, C) = (3, 4, 6)$  (i.e. all changes by T1 and T2 are written to disk)  
If T2 NOT successfully write to disk before the crash point  $\rightarrow (A, B, C) = (2, 1, 1)$  (i.e. only changes by T1 and written to disk)
- c) **Deadlock:** the situation when each set of transactions is waiting for a resource (e.g. lock) currently held by another transaction in the set. None can make progress. Example:  
**T1:**  $I_2(A)I_2(B)w_1(B) \dots$   
**T2:**  $I_2(B)I_2(A)w_2(A) \dots$   
T1 holds lock on A and request for lock on B, while T2 holds lock on B and request for lock on A. None of them releases the needed lock  $\rightarrow$  Keep waiting
- d) **Undo logging**  
Requires that data be written to disk immediately after a transaction finishes  
Increase number of disk I/Os  
Undo all uncommitted transactions, bottom-up (or latest-first)
- Redo Logging**  
Requires to keep all modified blocks in buffers until the transaction commits and the log records have been flushed  
Increases average number of buffers needed by transactions  
Redo all committed transaction, top-down (or earliest-first)

--End of Answers--

Solver: Phan Cong Minh